




Unravel Complexity. Connect Thought.

The cursor-native visual intelligence layer for desktop knowledge work. Instant spatial comprehension without leaving your active window.

The screenshot shows the Sutrr AI interface with a dark theme. At the top, there's a header with "Sutrr AI", a "Flowchart (98%)" indicator, a "100" score, and user avatars for Gemini, Groq, and Gemma. Below the header is a navigation bar with tabs: Summary, Auto, Flowchart, Timeline, Graph, Tree, and Mind Map. The main area displays a flowchart for a 2D matrix search algorithm. The flowchart starts with a green oval "Start: i=0, j=n-1", followed by a yellow decision diamond "i < m and j >= 0?". If "Yes", it goes to another yellow decision diamond "Matrix[i][j] == target?". If "Yes", it leads to a red oval "Return True". If "No" to either decision, it leads to a red oval "Return False". If "No" to the second decision, it leads to a blue rectangle "j -= 1", then a blue rectangle "i += 1", and loops back to the first decision diamond. The interface includes zoom controls (+, -, Fit, U) and a bottom section with "AI Enhanced" status, "Enriched by GEMMA", and a chat input field "Ask anything about this text...". At the very bottom, it says "This code implements an efficient search algorithm for a 2D matrix where rows and c..." and "98% Match • GEMMA".

The Invisible Cost of Context Switching

Knowledge workers waste hours daily translating dense information instead of executing.



FRICTION • CONTEXT SWITCH

4.2 hrs/wk

Lost in window juggling & re-reading

Users constantly Alt-Tab between IDEs, research PDFs, and browser AI chat boxes. Each window switch shatters deep cognitive focus and takes 15+ minutes to recover flow.

⚠ **Result: Cognitive exhaustion and fragmented output.**

DENSITY • COGNITIVE OVERLOAD

80% of Knowledge

Trapped in opaque, linear prose

Network protocols, smart contract traces, mathematical formulas, and 40-page NDAs are non-linear systems—yet forced into flat paragraphs that human brains struggle to absorb.

⚠ **Result: Information overload; inability to see structure.**

SECURITY • IP LEAKAGE

72% of Orgs

Restrict public cloud AI copy-paste

Workers paste confidential proprietary code, trade secrets, clinical trial findings, and private client agreements directly into external cloud LLM web boxes, risking severe regulatory breach.

⚠ **Result: Critical enterprise compliance barriers.**

Sutrr AI: Ambient Cursor-Native Intelligence



Never leave your flow. Highlight text anywhere, understand it instantly beside your cursor.

THE LEGACY WORKFLOW ✖

- 1. Highlighting friction** — Select text, copy to clipboard, Alt-Tab to browser.
- 2. Tab clutter & latency** — Open ChatGPT / Claude, paste prompt, wait 8 seconds.
- 3. Flat prose return** — Scan through another wall of generic bullet points.
- 4. Context destruction** — Alt-Tab back to work; cognitive momentum lost forever.

Result: High latency, fragmented focus, and cloud IP exposure.

THE SUTRR EXPERIENCE ✨

- 1. One Global Keystroke** — Press Ctrl+Shift+Space directly over any application.
- 2. Instant In-Situ HUD** — Native OS hook captures volatile RAM text in <50ms.
- 3. Spatial Visual Lenses** — Transforms words into directed flowcharts, ELI15, & analogies.
- 4. 100% Local & Private** — Zero cloud roundtrips; disappears the second HUD closes.

Result: Zero disruption. Muscle memory preserved. Instant spatial clarity.

Frictionless in 3 Simple Steps

Engineered to slot seamlessly into existing desktop muscle memory.



PHASE 01

Select Anywhere

Highlight text in Chrome, Adobe PDF, VS Code, Slack, Word, or Terminal. No browser plugin or per-app configuration required.

Supported: 100% of desktop software

PHASE 02

Trigger Global Hotkey

Press Ctrl+Shift+Space (or Cmd+Shift+Space on Mac). Native C# Win32 automation or Swift AX hook captures the text in volatile RAM.

Latency: <50ms in-memory capture

PHASE 03

Pick Your Lens

A floating HUD unfolds beside your cursor. Select Visual Flowchart, Plain English, Metaphor, or Comprehension Quiz for instant deconstruction.

Output: Spatial, interactive & exportable

One Input. Different Ways to Grasp It.

Human minds don't learn from one-size-fits-all prose. Sutrr adapts to your cognitive mode.



FLOWCHART & LOGIC MAP

Transforms sequential logic, algorithms, and distributed handshakes into directed SVG node maps.

PLAIN ENGLISH (Explain Like I'm 15)

Cuts through bureaucratic legalese, obscure academic phrasing, and opaque formulas to deliver the essence.

REAL-WORLD ANALOGY (Physical Metaphors)

Anchors abstract mathematical spaces, quantum mechanics, and consensus into tangible daily objects.

RAPID QUIZ (Active Retention Diagnostic)

Probes your comprehension with a targeted active recall question to verify you actually absorbed the nuance.

COMPARE & CONTRAST (Trade-Off Metrics)

Heuristically parses opposing protocols, specifications, and theories into side-by-side structured comparison tables.

EXECUTIVE SUMMARY (Key Takeaways)

Compresses 40-page contracts, research preprints, and dense documentation into 3 high-signal, decision-ready points.

CHRONOLOGICAL TIMELINE

Extracts timestamps, years, and historical milestones from text to automatically assemble an ordered chronological progression.

MIND MAP & HIERARCHY

Clusters central topics, technical keywords, and sub-clauses into a radial, multi-level categorical knowledge tree.



Privacy by Architecture, Not Marketing

DATA ISOLATION

0 Bytes

Saved or logged to external cloud

Selections reside strictly in volatile RAM. No screenshots taken, no screen scraping daemons. The text disappears the microsecond the HUD closes.

 **Never used for AI training or telemetry.**

OFFLINE CAPABILITY

100% Offline

Airplane Mode Verified

Deterministic syntactic parsing engine runs locally on client CPU/GPU. Synthesizes flowcharts, steps, and timelines without an active internet connection.

 **Zero roundtrip cloud latency.**

NATIVE PERFORMANCE

BYO-Model

Local Ollama & Private Cloud Ready

Detects local Ollama daemons (localhost:11434) out-of-the-box to run Gemma directly on your GPU, or connects to private VPC endpoints with zero cloud data retention.

 **Enterprise Bring-Your-Own-LLM**

Customer Acquisition & Revenue Model



STUDENTS

Rs 0 / mo

Free Forever for Individuals

- Generous free tier with local-only processing
- Student verification via .edu email or SheerID-style verification
- Free/discounted Researcher plan for the first **6–12 months**
- Goal: Become the default visual learning tool for coding & research classes

PRO INDIVIDUAL

Rs 380/ mo

Or Rs 4560 billed annually

- Affordable Rs 380/**month** Researcher plan
- Generate diagrams, flowcharts, research methods & result tables
- Cross-check outputs using multiple AI models
- Annual plan with **15–20% discount** for better retention & cash flow

Target: Developers, Researchers, Execs

B2B REVENUE SEGMENT

Annual Team Contracts

- Seat-based pricing with optional usage-based overages
- Upgrade path: Researcher → Team → Enterprise
- Team workspaces, sharing & collaboration
- Admin dashboards and security/compliance controls
- Private cloud / local deployment options

Target: Tech firms, Legal, Healthcare, Defense

Production Software in User Hands



Not a speculative pitch—a functional, engineered desktop reality.

Windows Setup + macOS Apple Silicon

Dual-platform production releases live. Windows installer with automatic system tray & uninstallation; zero-install standalone portable .exe; and native universal macOS .dmg.

Status: Signed, tested, and actively downloadable.

Universal Desktop Compatibility

Operates across 100% of active desktop applications—Chrome, Adobe Acrobat, VS Code, Slack, Microsoft Office, and Terminals—without requiring per-app plugins or browser extensions.

Status: Signed, packaged, and actively downloadable.

8 Multi-Modal Understanding Engines Built

Complete heuristic and NLP rule engine supporting Directed Flowcharts, Plain-English Deconstruction, Analogies, Quizzes, Comparison Matrices, and Dynamic Simulations.

Status: Full interactive branching suite operational.

Hybrid Local-First & Multi-Model Engine

Deterministic local spaCy NLP engine coupled with an agentic LangGraph pipeline featuring automatic schema validation and auto-correction; seamlessly bridges to local Ollama (Gemma) or cloud APIs (Gemini/Groq).

Status: Resilient multi-provider routing live.